



HPC Software Verification in Action: A Case Study with Tensor Transposition

Erdal Mutlu, Ajay Panyala, Sriram Krishnamoorthy

Post Doctoral Research Associate

erdal.mutlu@pnnl.gov



PNNL is operated by Battelle for the U.S. Department of Energy



Correctness in HPC

- Verification challenges in HPC
 - hard to specify scientific applications
 - different sources of non-determinism
 - ✓ concurrency
 - ✓ approximation
 - ✓ randomness
 - floating point arithmetic
- Where to start?
 - functional equivalence checking
 - ✓ optimized implementation vs simple trusted implementation
 - ✓ model checking techniques
 - ✓ symbolic execution

Tensor Algebra

- Tensor Operations
 - majority of machine learning (ML) application
 - scientific applications: quantum chemistry
- Tensor Transposition

$$A_{\Pi(i_0, i_1, \dots, i_{d-1})} \leftarrow B_{i_0, i_1, \dots, i_{d-1}}$$

- E.g. 2-D Matrix transposition : $A(j, i) = B(i, j)$

```
for(int i = 0; i < row_size; i++){
    for(int j = 0; j < col_size; j++){
        src_id = i * col_size + j;
        dest_id = j * row_size + i;
        A[dest_id] = B[src_id];
    }
}
```

Tensor Transpose Implementations

- NWChem implementation
 - specialized implementations for 2-D and 4-D tensors
 - optimized for production use
 - uses heavy pointer arithmetic
- TAMM implementation
 - generalized C++ implementation
 - supports reduction
- Reference implementation
 - specialized for each permutation

```
int void ref_impl_3D_0_2_1(double *dest,
    in                           const double *src,
void tce_sort(double *A, double *B, /*...*/) {
    //...
    for (i = 0; i < bSize[3]; i++) {
        aPtr[2] = aPtr[3];
        for (j = 0; j < bSize[2]; j++) {
            aPtr[1] = aPtr[2];
            for (k = 0; k < bSize[1]; k++) {
                aPtr[0] = aPtr[1];
                for (l = 0; l < bSize[0]; l++) {
                    *bPtr++ = *aPtr[0];
                    aPtr[0] += bJump[0];
                }
                aPtr[1] += bJump[1];
            }
            aPtr[2] += bJump[2];
        }
        aPtr[3] += bJump[3];
    }
    //...
}
    } else {
        ref_impl_3D_0_1_2(dest, src, dims);
    }
}
}
```

Equivalence Checking Tools

- CIVL: the concurrency intermediate verification language
 - symbolic execution-based model checking
 - supports many HPC-relevant parallel programming models
 - ✓ OpenMP, pthreads, CUDA
 - checks different properties
 - ✓ race freedom, deadlock freedom, assertion violation
- CodeThorn
 - trace-based equality checking technique
 - re-write updates into a normalized form
 - compare resulting traces

Research Questions

- Can current equivalence checking tools be used in **production code**?
- How **scalable** are these tools' performance?
- Where these tools **fall short**?

Experimental Setup

- Compared 2, 3 and 4 dimensional tensor transpose implementations
 - NWChem
 - TAMM
 - Reference
- CIVL setup
 - Bounded input size
 - ✓ `$assume(1 <= N <= BOUND) ;`
 - Fixed input size
 - ✓ `$input N = X;`
- CodeThorn setup
 - Fixed input size

Research Questions

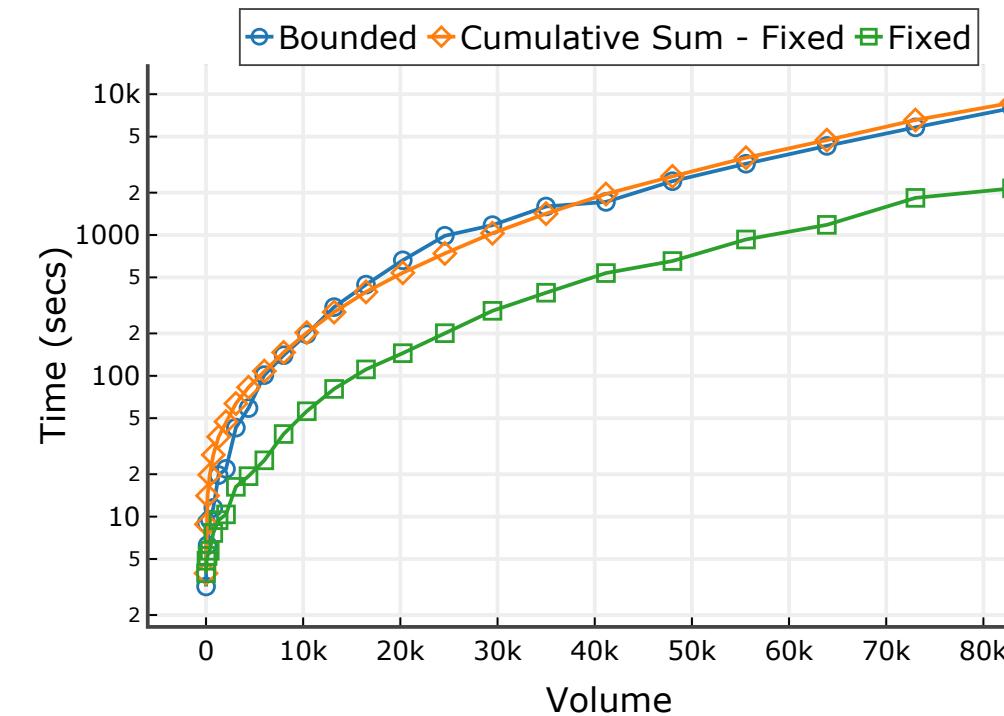
- Can current equivalence checking tools be used in **production code**?
 - Worked for us (with some code modification) ☺
 - Valuable input during development
- How **scalable** are these tools?
- Where these tools **fall short**?

Research Questions

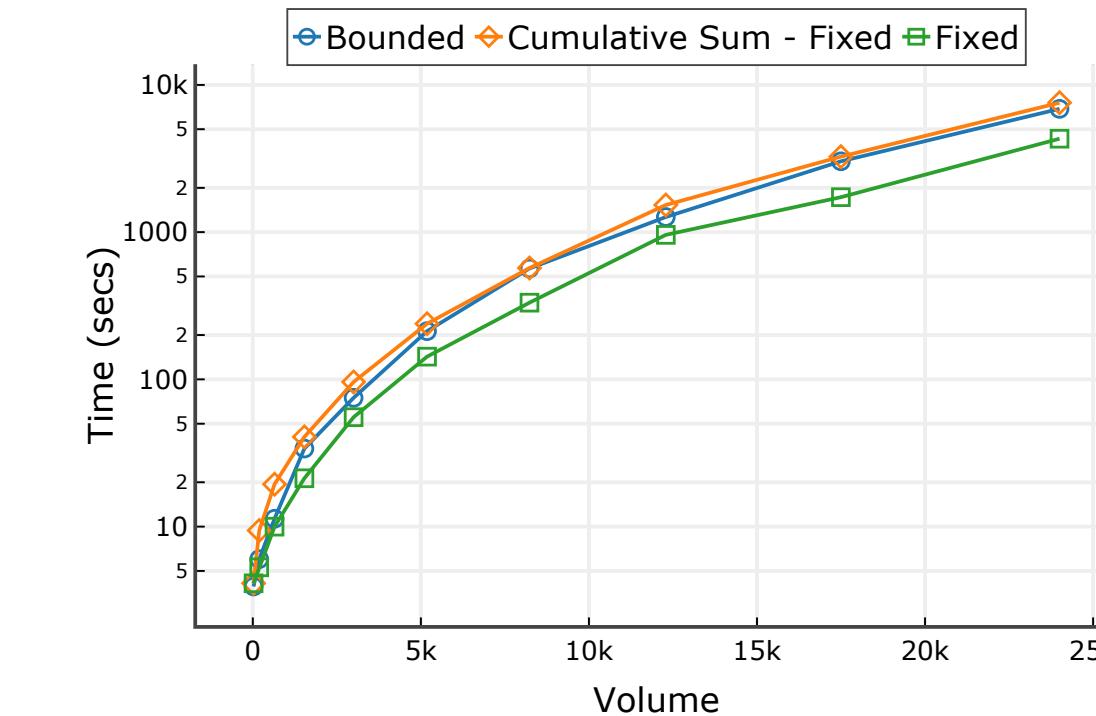
- Can current equivalence checking tools be used in **production code**?
 - Worked for us (with some code modification) ☺
 - Valuable input during development
- How **scalable** are these tools?
 - Works relatively well in small size inputs
 - Using different modes helps
- Where these tools fall **short**?

Results – CIVL

- Functional equality checking for NWChem against reference implementation



3D - NWChem vs Ref

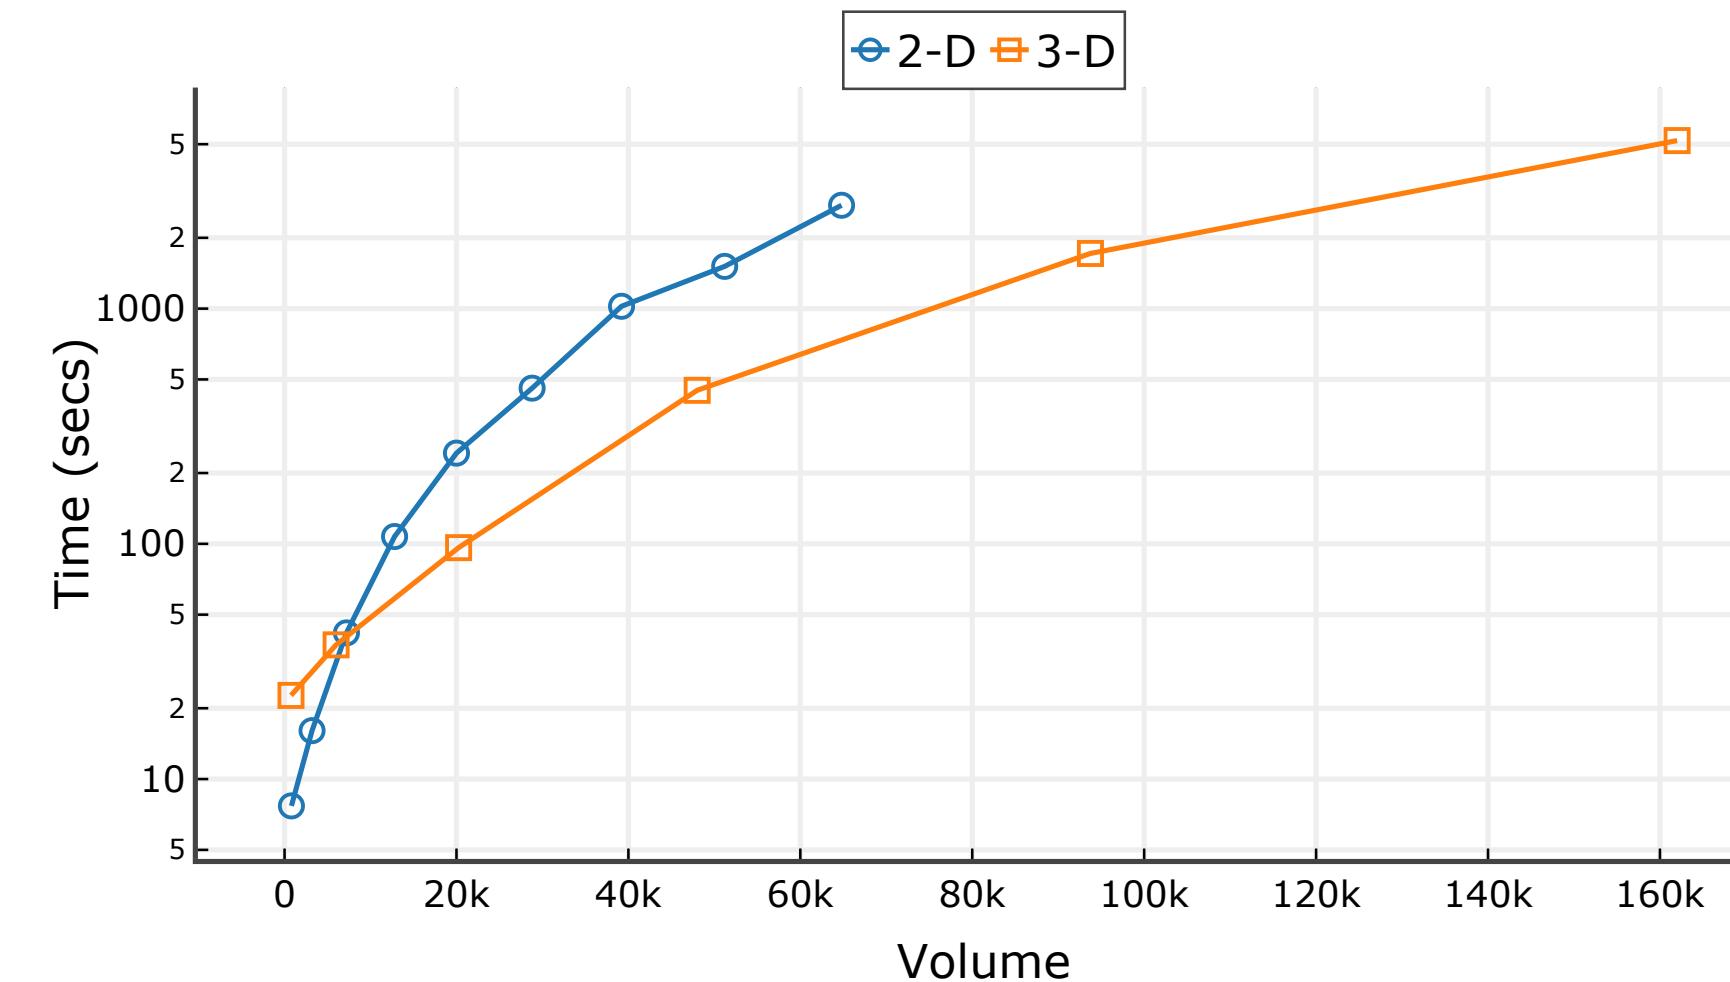


4D - NWChem vs Ref

Volume = # of Permutations * Loop Volume

Results – CodeThorn

- Functional equality checking with CodeThorn



Research Questions

- Can current equivalence checking tools be used in **production code**?
 - Worked for us (with some code modification) ☺
 - Valuable input during development
- How **scalable** are these tools?
 - Works relatively well in small size inputs
 - Using different modes helps
- Where these tools **fall short**?
 - Supporting programming languages (C++)
 - Problems with pointer arithmetic

Conclusion

- Emerging heterogeneous architectures in HPC:
 - complex optimized versions of the same algorithm
 - hard to write specify/verify
 - floating point arithmetic
- Functional equivalence checking
 - both, CIVL and CodeThorn, can be used in production codes
 - symbolic evaluation based
 - works relatively well for small input size
 - some caveats:
 - ✓ subset of C programming language is supported
 - ✓ false positives (case with pointer arithmetic)



Pacific
Northwest
NATIONAL LABORATORY

Thank you



False Positive – CIVL

```
void tce_sort(double *A, double *B, /*...*/) {
    //...
    for (i = 0; i < bSize[3]; i++) {
        aPtr[2] = aPtr[3];
        for (j = 0; j < bSize[2]; j++) {
            aPtr[1] = aPtr[2];
            for (k = 0; k < bSize[1]; k++) {
                aPtr[0] = aPtr[1];
                for (l = 0; l < bSize[0]; l++) {
                    *bPtr++ = *aPtr[0];
                    aPtr[0] += bJump[0];
                }
                aPtr[1] += bJump[1];
            }
            aPtr[2] += bJump[2];
        }
        aPtr[3] += bJump[3];
    }
    //...
}
```

