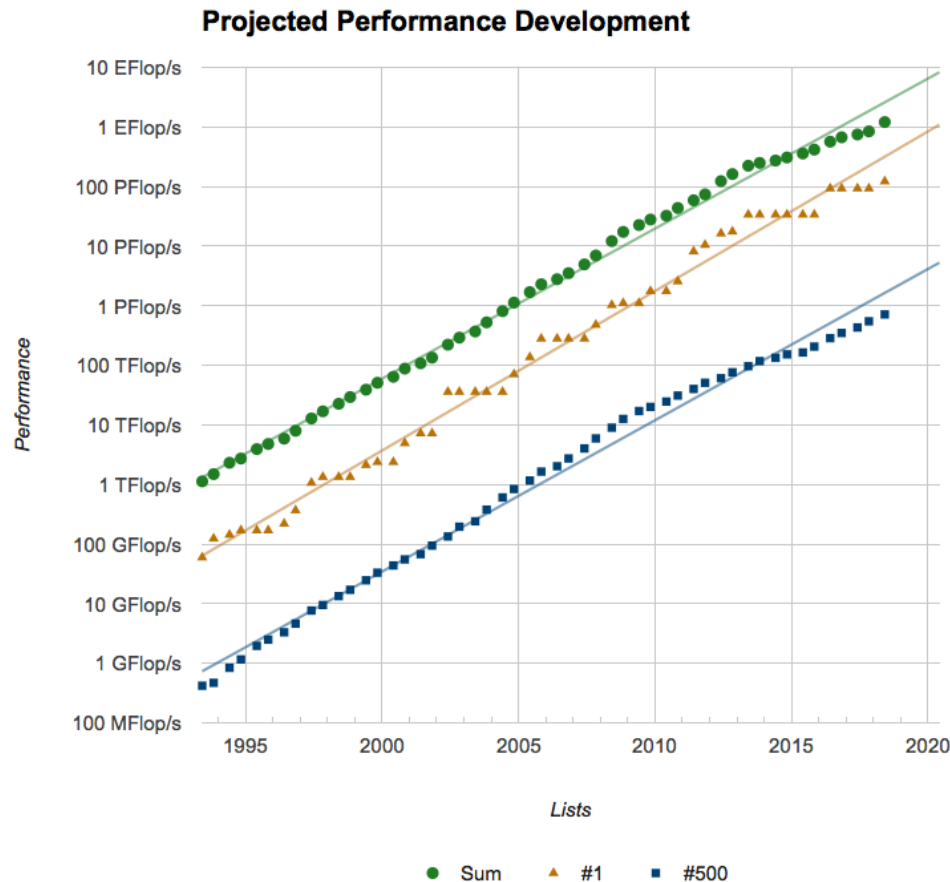


PARCOACH Extension for a Full-Interprocedural Collectives Verification

Pierre Huchant^{1,2}, **Emmanuelle Saillard**², Denis Barthou^{1,2},
Hugo Brunie^{1,2,3} and Patrick Carribault³

1. Univ. of Bordeaux, Bordeaux INP - 2. Inria Bordeaux - 3. CEA, DAM, DIF, F-91297 Arpajon, France

Correctness 2018
Dallas, Texas

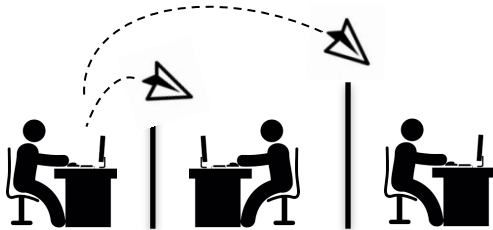


- More complex applications (i.e., combination of parallel programming models)
- Machines more complex, heterogeneous architectures
- Exascale systems targeted in 2020

How can we help developers having correct parallel applications?

⇒ Precision ⇒ Scalability ⇒ Soundness ⇒ Heterogeneity

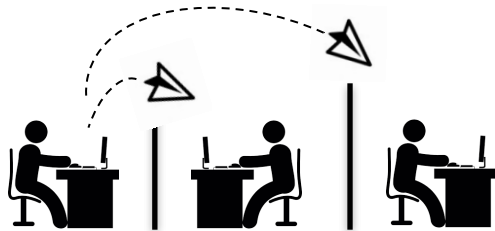
MPI



OpenMP



MPI



Blocking or non-blocking communication involving all MPI processes of the same communicator

`MPI_Barrier, MPI_Ibarrier, MPI_Bcast, ...`

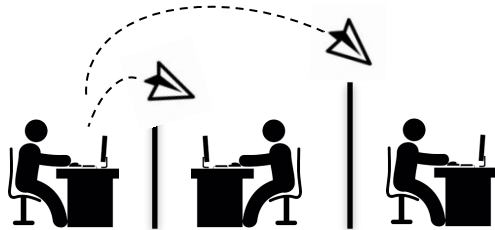
OpenMP



Barrier and worksharing construct

`#pragma omp barrier/single/for/
sections/workshare`

MPI



Blocking or non-blocking communication involving all MPI processes of the same communicator

`MPI_Barrier, MPI_Ibarrier, MPI_Bcast, ...`

OpenMP

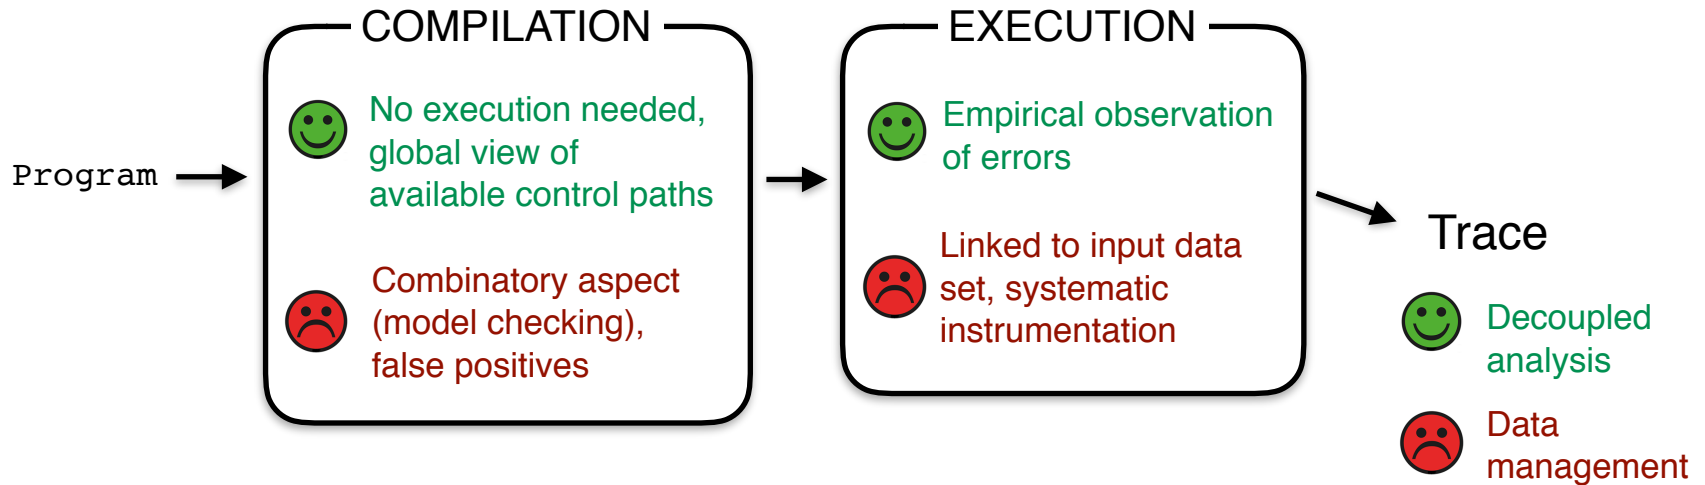


Barrier and worksharing construct

`#pragma omp barrier/single/for/
sections/workshare`

MPI/OpenMP specification : All MPI processes / OpenMP threads must have the same sequence of collectives

Goal : Detect collective errors (i.e., collective mismatch)



- **Static tools**

- > MPI-SPIN, TASS, OAT, Zhang *et al.*, GCC, ICC, LLVM

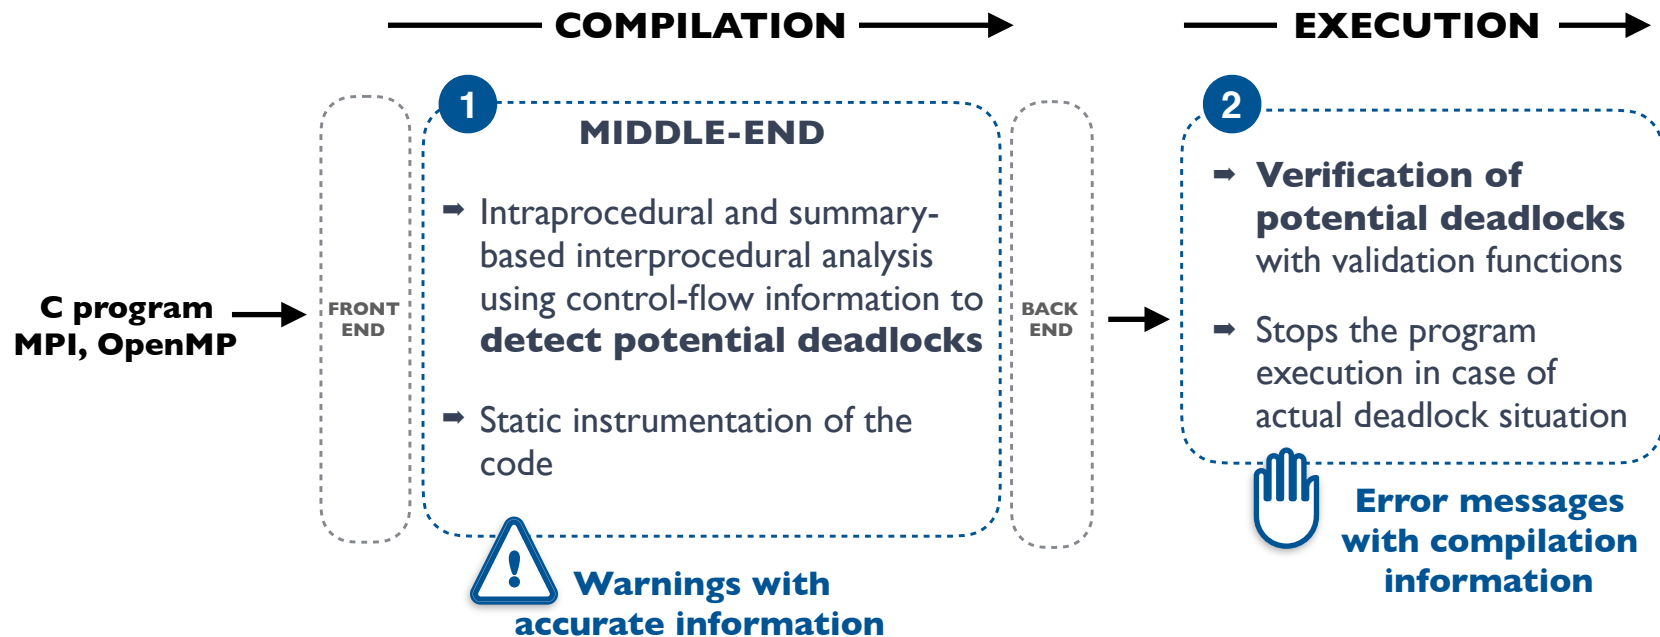
- **Dynamic tools**

- > MUST, DAMPI, MPI profiling library (MPICH2, MPI/SX)

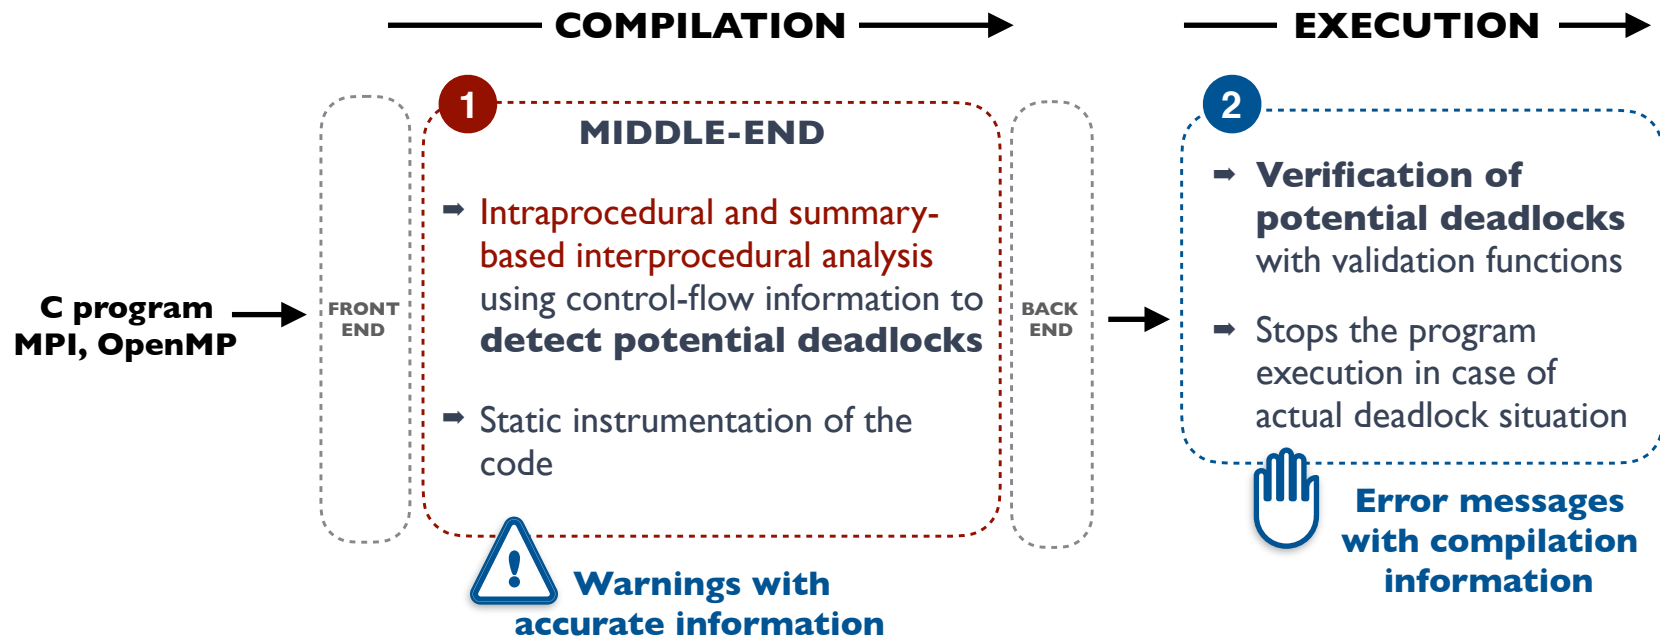
- **Trace-based dynamic tools**

- > IMC, Intel Inspector XE, Sun thread analyzer

PARallel COntrol flow Anomaly CHecker



PARallel COntrol flow Anomaly CHecker



Contributions

- More accurate interprocedural analysis
- Full integration into the LLVM compiler


```
void c( ) {  
  
    if( )  
        MPI_Barrier(com2);  
    else  
        /*...*/  
}  
  
int main( ) {  
  
    /*...*/  
    MPI_Barrier(com1);  
  
    if( )  
        c( );  
  
    /*...*/  
  
    MPI_Finalize();  
  
}
```



```
void c( ) {  
  
    if( )  
        MPI_Barrier(com2);  
    else  
        /*...*/  
}  
  
int main( ) {  
  
    /*...*/  
    MPI_Barrier(com1);  
  
    if( )  
        c( );  
  
    /*...*/  
  
    MPI_Finalize();  
  
}
```



```
void c( ) {  
  
    if( )  
        MPI_Barrier(com2);  
    else  
        /*...*/  
}  
  
int main( ) {  
  
    /*...*/  
    MPI_Barrier(com1);  
  
    → if( )  
        c( );  
    /*...*/  
    MPI_Finalize();  
}
```

```
void c( ) {  
  
    if( )  
        MPI_Barrier(com2);  
    else  
        /*...*/  
}  
  
int main( ) {  
  
    /*...*/  
    MPI_Barrier(com1);  
  
    if( )  
        c( );  
        /*...*/  
    MPI_Finalize();  
}
```



```
void c( ) {
```

→

```
    if( )
```

```
        MPI_Barrier(com2);
```

```
    else
```

```
        /*...*/
```

```
}
```

```
int main( ) {
```

```
    /*...*/
```

```
    MPI_Barrier(com1);
```

```
    if( )
```

```
        c( );
```

→

```
    /*...*/
```

```
    MPI_Finalize();
```

```
}
```

```
void c( ) {
```

```
    if( )
```

```
        MPI_Barrier(com2);
```

```
    else
```

```
        /*...*/
```

```
}
```

```
int main( ) {
```

```
    /*...*/
```

```
    MPI_Barrier(com1);
```

```
    if( )
```

```
        c( );
```

```
    /*...*/
```

```
    MPI_Finalize();
```

```
}
```

```
void c( ) {
```

```
    if( )
```

```
        MPI_Barrier(com2);
```

```
    else
```

```
        /*...*/
```

```
}
```

```
int main( ) {
```

```
    /*...*/
```

```
    MPI_Barrier(com1);
```

```
    if( )
```

```
        c( );
```

```
    /*...*/
```

```
    MPI_Finalize();
```

```
}
```

Deadlock

The machine state does not help to detect the source of the deadlock

```
void c( ) {
```

```
    if( )
```

```
        MPI_Barrier(com2);
```

```
    else
```

```
        /*...*/
```

```
}
```

```
int main( ) {
```

```
    /*...*/
```

```
    MPI_Barrier(com1);
```

```
    if( )
```

```
        c( );
```

```
    /*...*/
```

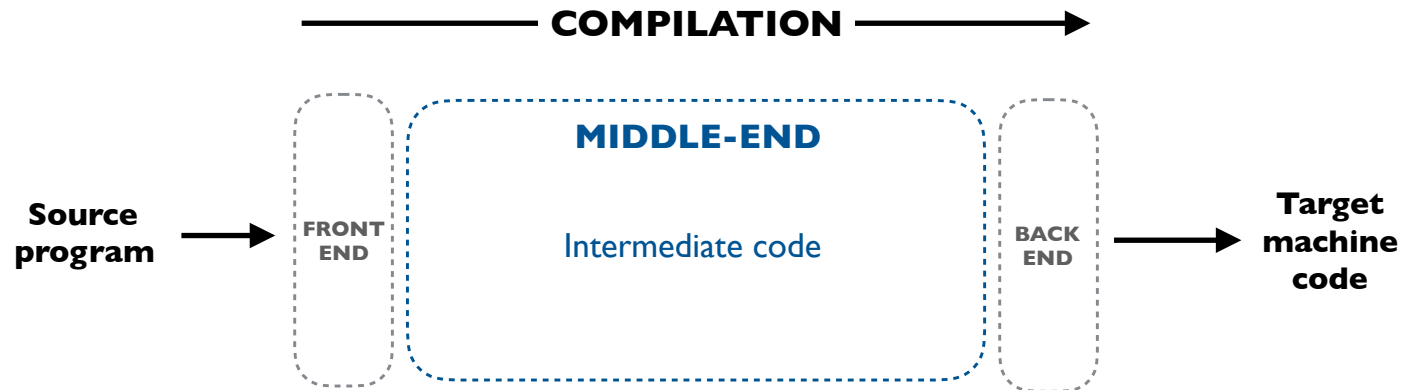
```
    MPI_Finalize();
```

```
}
```

Deadlock

The machine state does not help to detect the source of the deadlock

Structure of a compiler



Intermediate representation

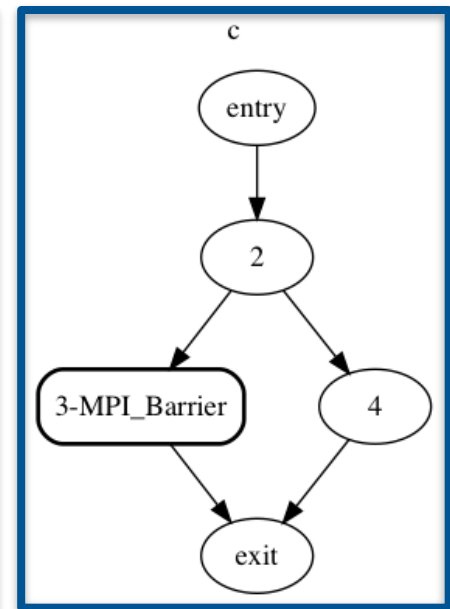
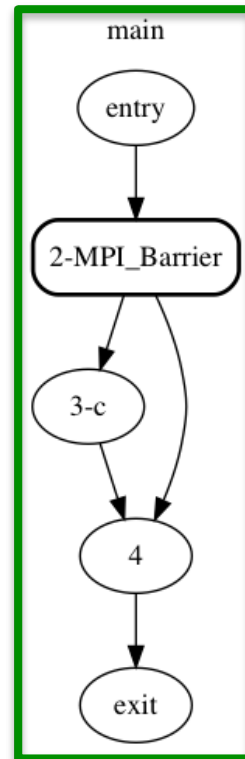
- Programming language independent
- Target machine independent

Control Flow Graph (CFG)

- Models all program executions
- Right representation to study instruction flow

Control Flow Graph

```
void c( ) {  
  
    if( ) } 2  
        MPI_Barrier(com2); } 3  
    else  
        /*...*/ } 4  
}  
  
int main( ) {  
  
    /*...*/  
    MPI_Barrier(com1); } 2  
  
    if( )  
        c( ); } 3  
    /*...*/  
    MPI_Finalize(); } 4  
}
```

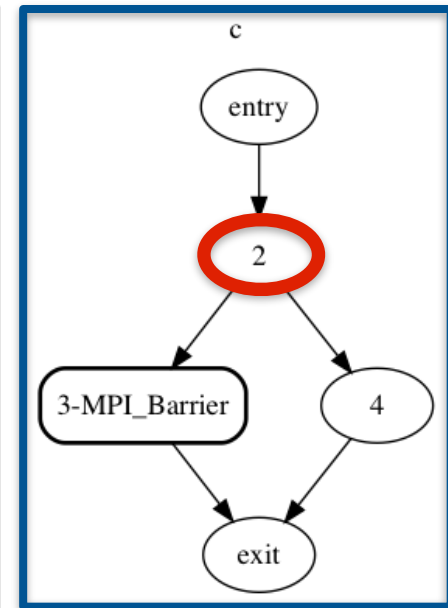
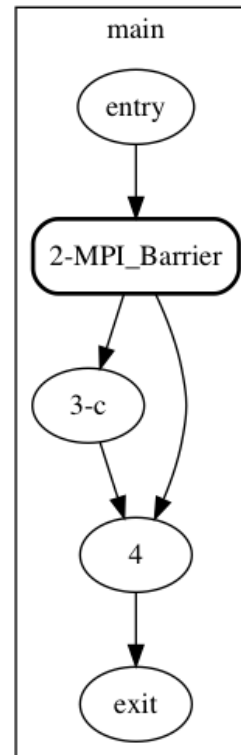


Iterated post dominance frontier (PDF⁺)

U : Set of nodes in PPCFG

PDF⁺(U) : Set of nodes that can lead both to a node in U or not

$$\text{PDF}^+(\{3_c\}) = \{2_c\}$$



Summary of function c: \emptyset

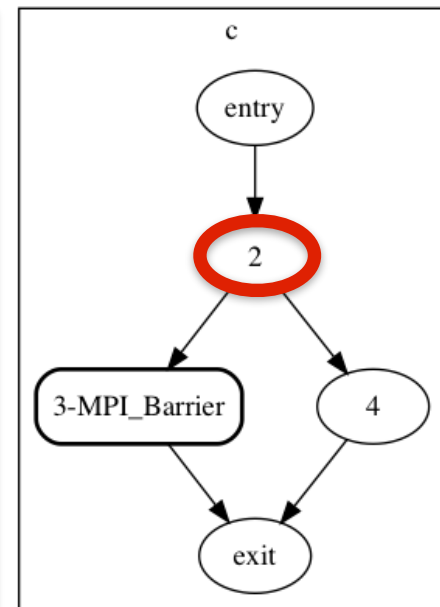
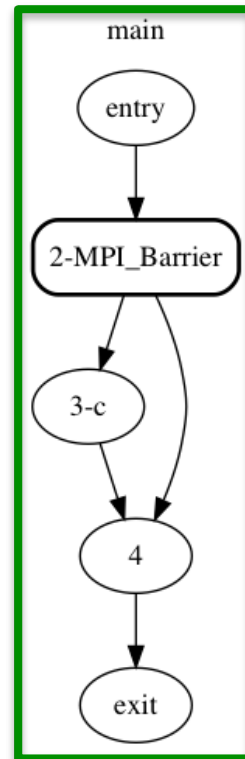
Iterated post dominance frontier (PDF+)

U : Set of nodes in PPCFG

PDF+(U) : Set of nodes that can lead both to a node in U or not

$$\text{PDF}^+(\{2_{\text{main}}\}) = \emptyset$$

$$\text{PDF}^+(\{3_c\}) = \{2_c\}$$

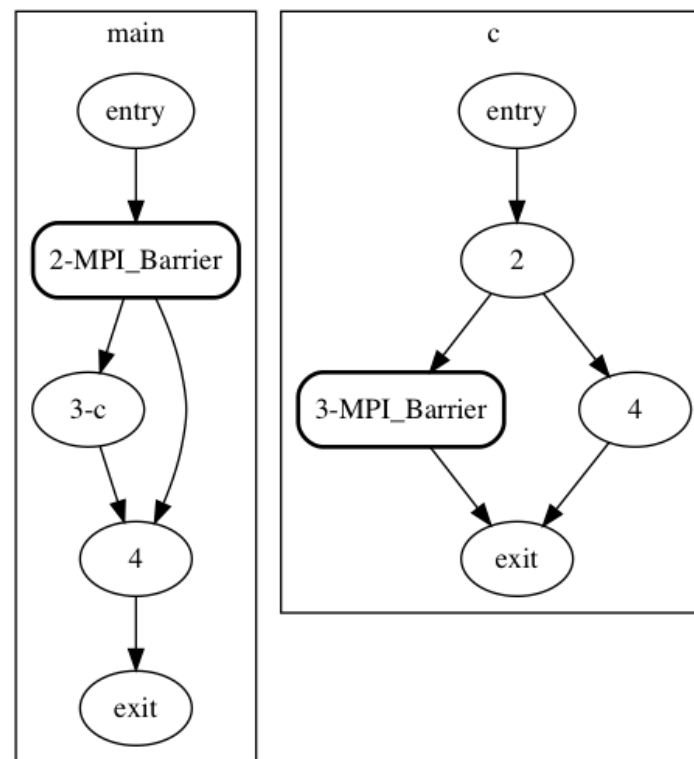


Summary of function c: \emptyset

Summary-based interprocedural analysis = Intraprocedural analysis

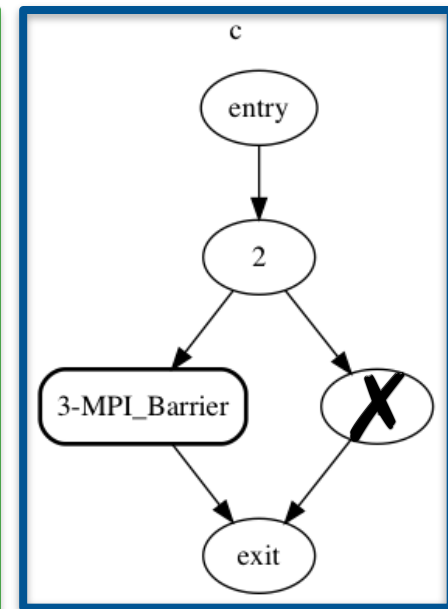
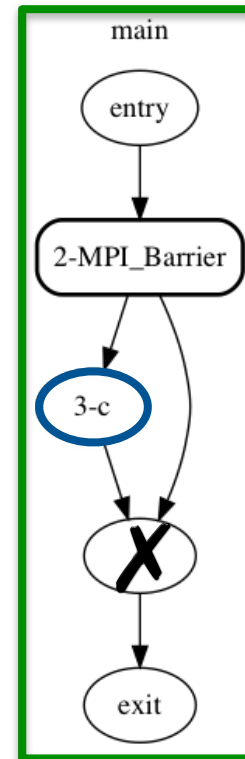
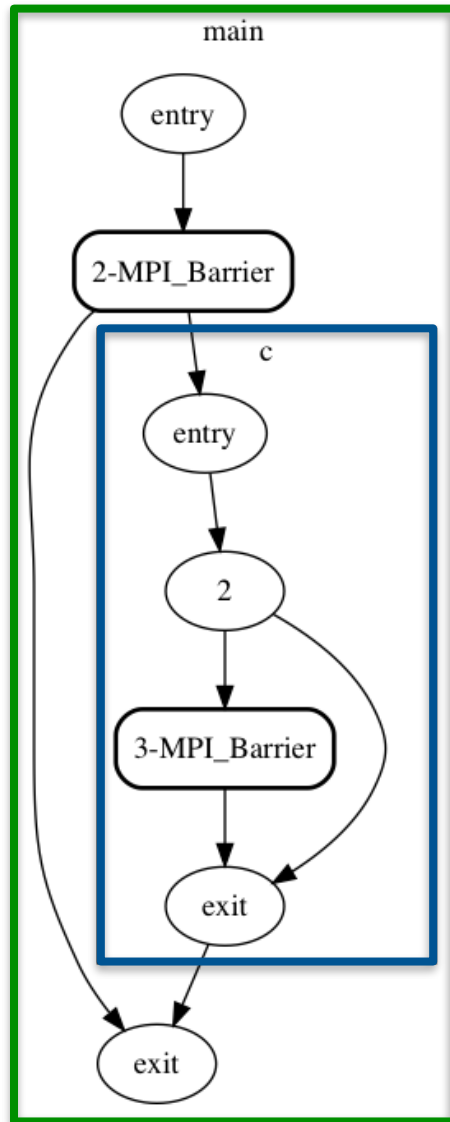


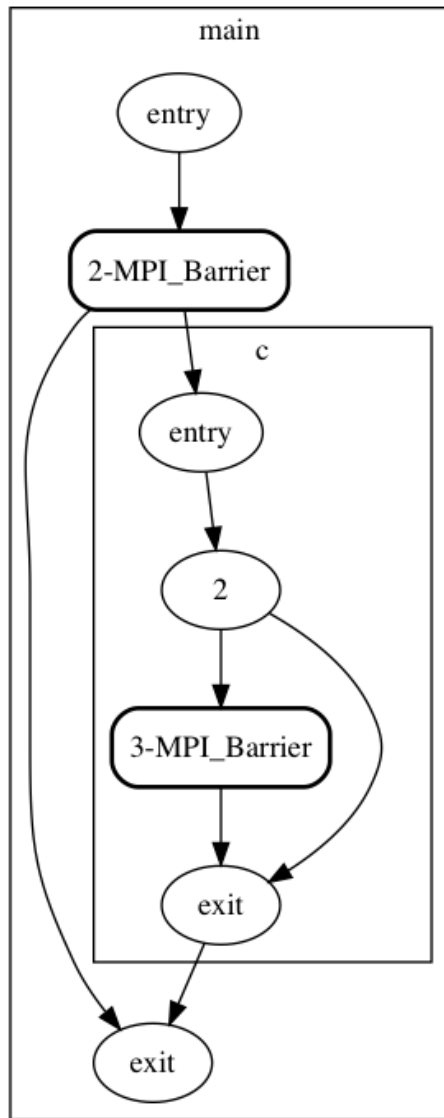
Parallel Program Control Flow Graph





Parallel Program Control Flow Graph

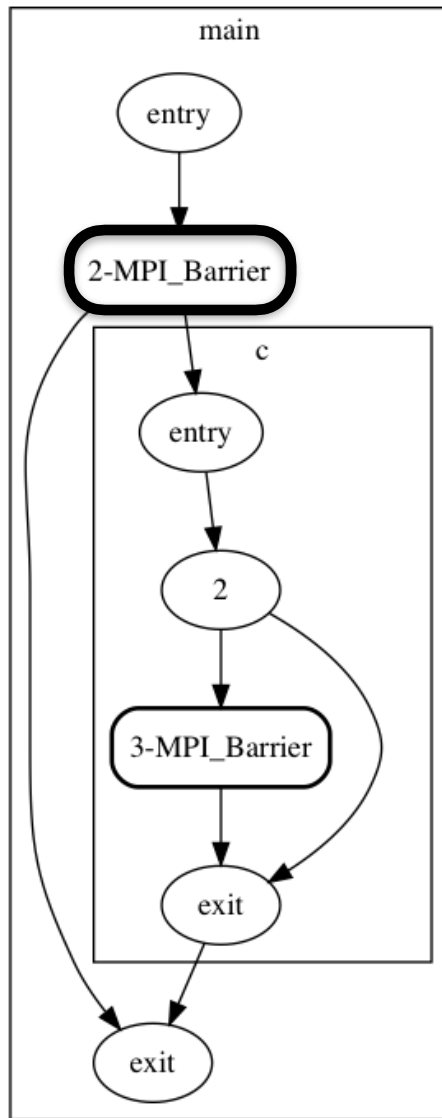




Iterated post dominance frontier (PDF⁺)

U : Set of nodes in PPCFG

PDF⁺(U) : Set of nodes that can lead both to a node in U or not

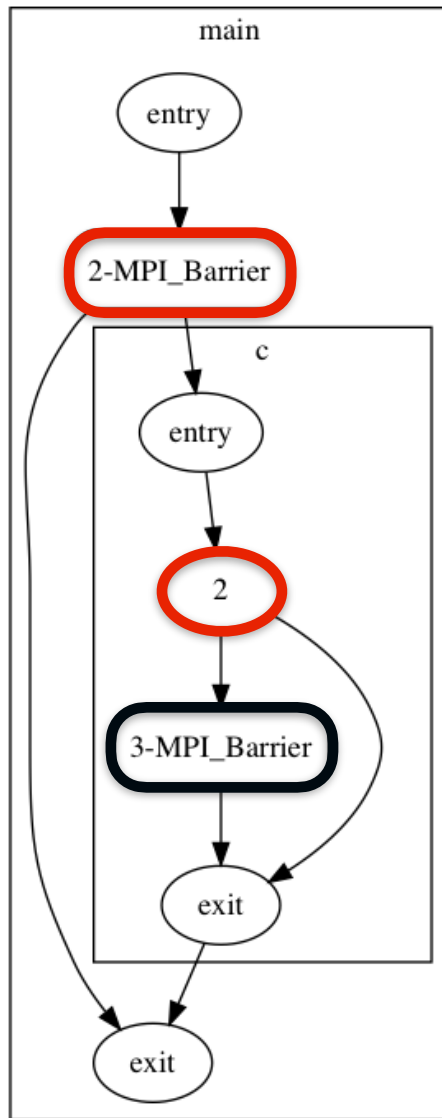


Iterated post dominance frontier (PDF⁺)

U : Set of nodes in PPCFG

PDF⁺(U) : Set of nodes that can lead both to a node in U or not

$$\text{PDF}^+(\{2_{\text{main}}\}) = \emptyset$$



Iterated post dominance frontier (PDF⁺)

U : Set of nodes in PPCFG

PDF⁺(U) : Set of nodes that can lead both to a node in U or not

$$\text{PDF}^+(\{2_{\text{main}}\}) = \emptyset$$

$$\text{PDF}^+(\{3_c\}) = \{2_c, 2_{\text{main}}\}$$

```
1 void c( ) {  
2  
3  if( )  
4    MPI_Barrier(com2);  
5  else  
6    /*...*/  
7  }  
8  
9  int main( ) {  
10  
11    /*...*/  
12    MPI_Barrier(com1);  
13  
14    if( )  
15      c( );  
16  
17    MPI_Finalize();  
18  }
```

What a user can read on stderr

PARCOACH: **warning:** MPI_Barrier line 4
possibly not called by all processes
because of conditional(s) line(s) 3, 14

```
void c( ) {  
  
    if( )  
  
        MPI_Barrier(com2);  
    else  
        /*...*/  
}  
  
int main( ) {  
  
    /*...*/  
    MPI_Barrier(com1);  
  
    if( )  
        c( );  
  
    /*...*/  
  
    MPI_Finalize();  
}
```

- No warning = no instrumentation

```
void c( ) {

    if( )
        CC(MPI, com2, i_barrier, 0)
        MPI_Barrier(com2) X
    else
        /*...*/
}

int main( ) {

    /*...*/
    MPI_Barrier(com1); ✓

    if( )
        c( );

    /*...*/

    CC(MPI, com2, 0, Ø )

    MPI_Finalize();

}
```

- No warning = no instrumentation

- If at least one warning:

> Insert a Check Collective (CC) function

CC(i_{model}, com_c, i_c, 0) before each collective c, **starting with the 1st collective that may deadlock**

> Insert **CC(i_{model}, com, 0, Ø)** before exit statements and some MPI functions (MPI_Abort, MPI_Finalize)

* i_{model} = Parallel programming model used

* com_c = communicator related to c (0 for OpenMP)

* i_c = collective ID

* Ø = collectives that may deadlock (set generated at compile-time)



- GCC plugin



LLVM Pass

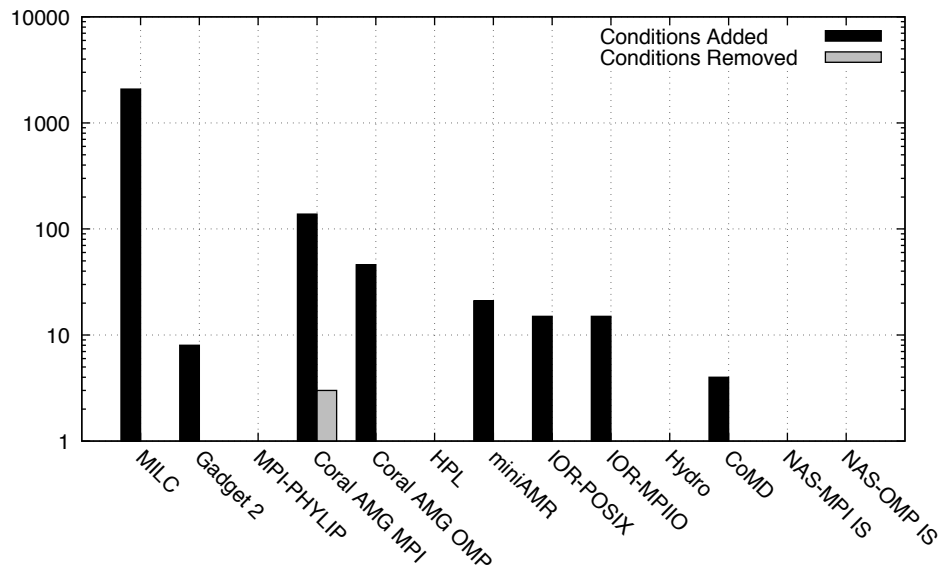


- Intraprocedural Analysis VS Full-interprocedural Analysis
- Applications and Benchmarks Statistics

Application	Parallelism	# func.	# coll.	# com.
MILC*	MPI	24,242	635	253
Gadget-2	MPI	193	70	1
MPI-PHYLIP*	MPI	4,000	128	12
Bench. / mini app.	Parallelism	# func.	# coll.	# com.
Coral AMG	MPI	1,207	79	19
	OpenMP	1,207	11	-
HPL	MPI	193	3	1
miniAMR	MPI	103	43	2
IOR-POSIX	MPI	175	82	5
IOR-MPIIO	MPI	197	88	5
Hydro	MPI	99	13	1
CoMD	MPI	124	8	1
NAS-MPI IS	MPI	36	9	1
NAS-OMP IS	OpenMP	51	3	-



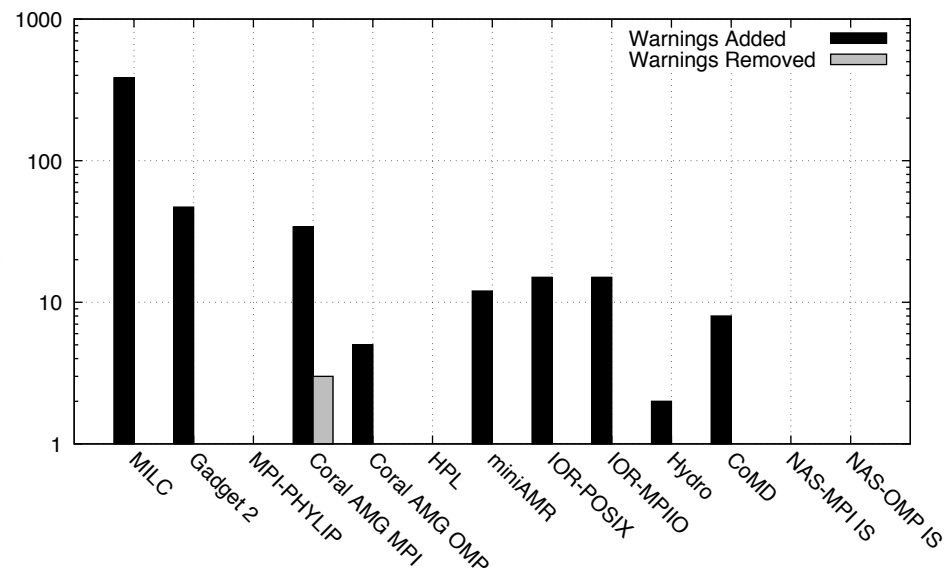
Experimental Results



Conditionals
Intra VS Full-inter

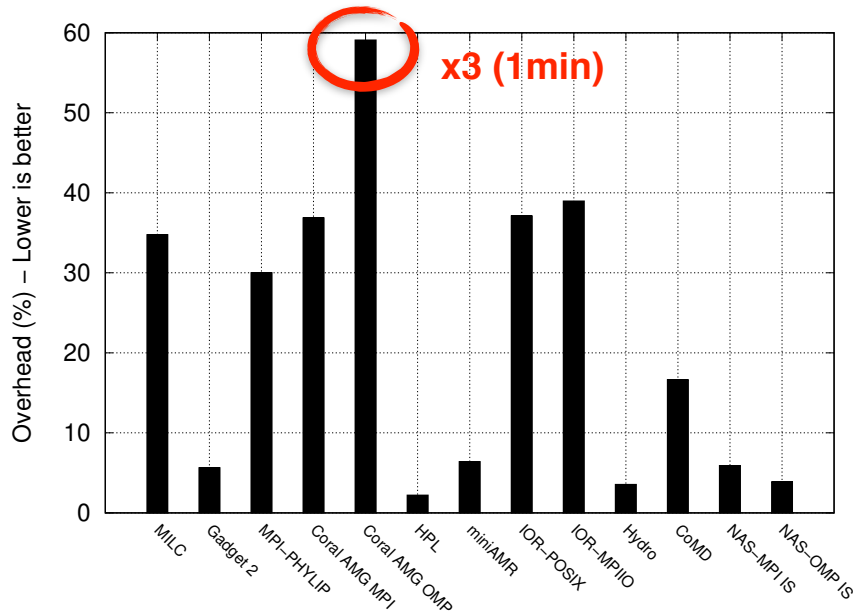
Warnings

Intra VS Full-inter





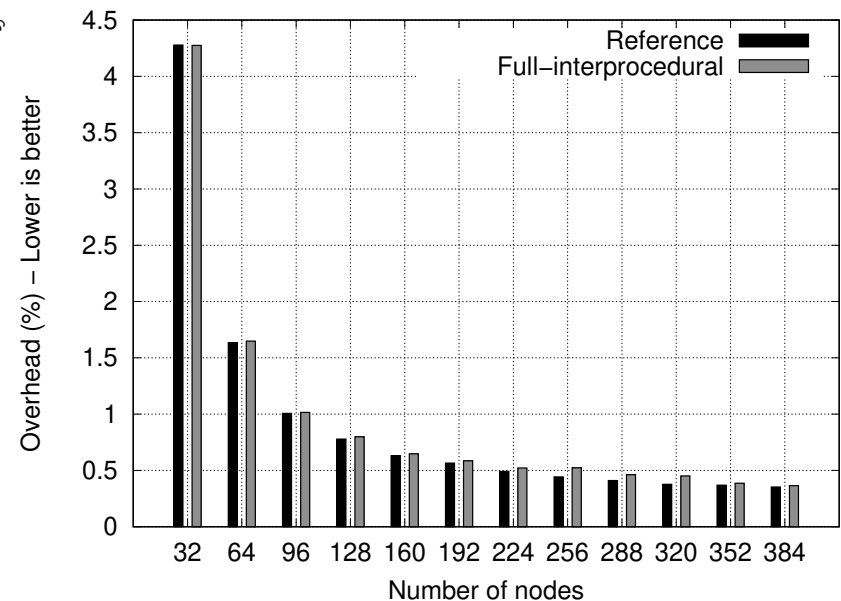
Overhead



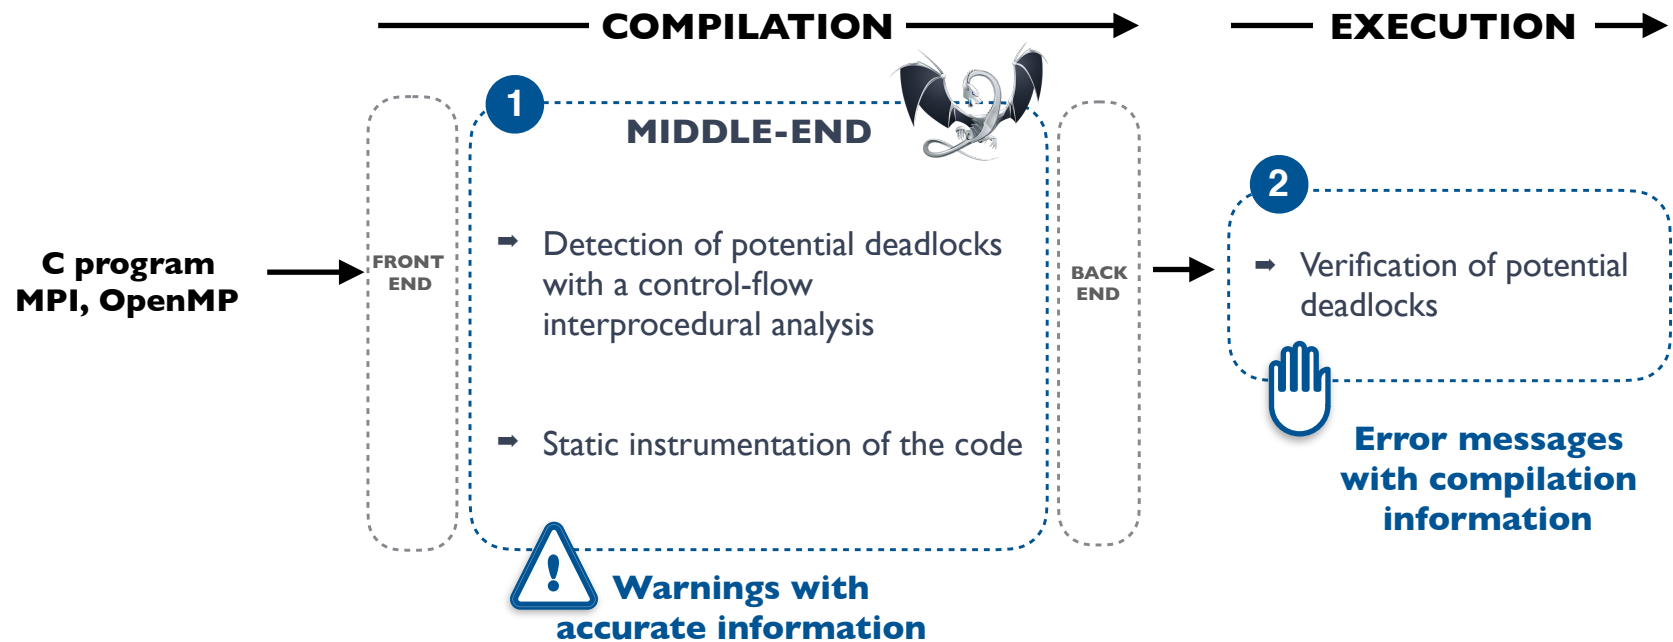
Execution-time overhead

Hydro benchmark

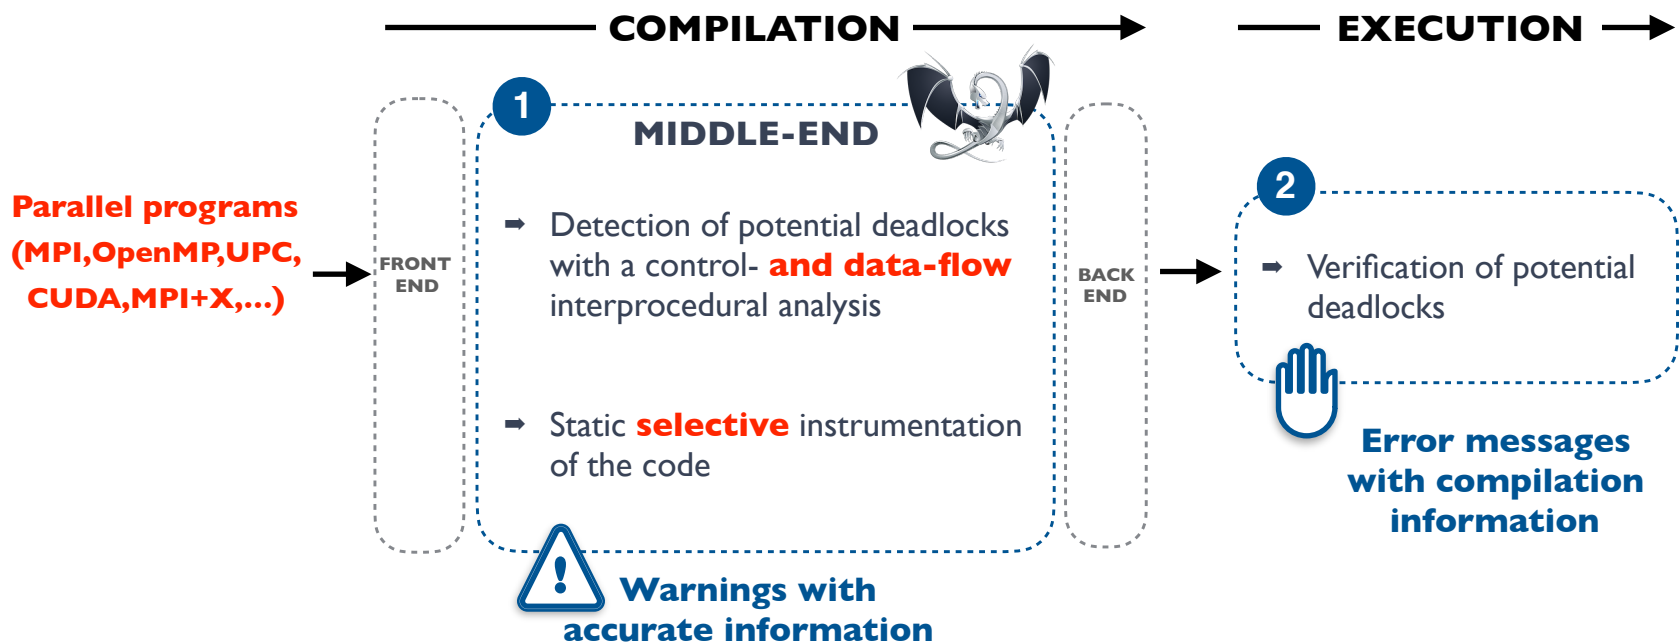
Compile-time overhead



PARallel COntrol flow Anomaly CHecker (PARCOACH)



PARallel COntrol flow Anomaly CHecker (PARCOACH)



➔ Precision ➔ Scalability ➔ Soundness ➔ Heterogeneity

Thank you!

Follow us on www.inria.fr